

# ShopSmartlytoday - Production-Ready Architecture

This document translates the provided storefront HTML/CSS into a scalable, secure, and observable e-commerce platform. It is opinionated for AWS (fully portable to Azure/GCP).

---

## 1) High-Level Overview

**Goal:** Serve a global, performant e-commerce experience with reliable search, secure checkout, and operational excellence.

**Architecture style:** Modular microservices with an API Gateway + event-driven backbone.

**Core tenets:** - **Performance first:** CDN, edge caching, image optimization, aggressive read caching. - **Security by default:** Zero-trust networking, managed auth, WAF, secrets vault, least privilege. - **Observability:** End-to-end tracing, structured logs, SLOs, synthetic tests. - **Automation:** IaC, CI/CD, canary/blue-green, automated rollbacks.

---

## 2) Reference Diagram (Mermaid)

```
graph TD
    subgraph Edge
        DNS[Route53 DNS]
        CDN[CloudFront CDN]
        WAF[AWS WAF]
    end

    subgraph Web
        FE[Next.js/React SPA+SSR]
        ASSET[S3 Static Assets]
    end

    subgraph Platform
        APIGW[API Gateway]
        AUTH[Cognito / Keycloak]
        BFF[Node.js BFF]
    end

    subgraph Services
        CAT[Catalog Service]
        SRCH[Search Service (OpenSearch)]
    end
```

```

CART[Cart Service]
ORD[Order Service]
PAY[Payments Adapter]
INV[Inventory Service]
REC[Recommendations]
UMS[User/Accounts]
CMS[CMS (admin)]
end

subgraph Data
PG[(PostgreSQL/Aurora)]
REDIS[(ElastiCache/Redis)]
OS[(OpenSearch)]
OBJ[(S3 Product Media)]
BUS[(Kafka/MSK)]
WH[(Redshift/Snowflake)]
end

subgraph Observability
LOGS[(CloudWatch/ELK)]
MET[Prometheus/Grafana]
TRACE[OpenTelemetry->Jaeger/X-Ray]
end

DNS-->CDN-->WAF-->FE
FE-->ASSET
FE-->APIGW
APIGW-->AUTH
APIGW-->BFF

BFF-->CAT
BFF-->SRCH
BFF-->CART
BFF-->ORD
BFF-->PAY
BFF-->INV
BFF-->UMS
BFF-->REC
BFF-->CMS

CAT-->PG
CART-->REDIS
ORD-->PG
INV-->PG
SRCH<-->OS
PAY-->|Stripe/Adyen|EXT[Payment PSP]
CAT-->OBJ

```

```
ORD--events-->BUS  
INV--events-->BUS  
PAY--events-->BUS  
BUS-->WH
```

```
Services-->LOGS  
Services-->MET  
Services-->TRACE
```

### 3) Mapping UI → Services

- **Header/Search bar** → **Search Service** (autocomplete, typeahead, filters), **Catalog** for categories.
- **Category tiles** → **Catalog Service** (taxonomy), **CMS** for promotional copy/images.
- **Featured Deals/Recommended** → **CMS** + **Recommendations** (rules + ML). Edge cached.
- **Cart/Orders** → **Cart Service** (ephemeral Redis) + **Order Service** (Aurora, transactional).
- **Auth/Account** → **UMS** + **Cognito/Keycloak** (OIDC, social login).
- **Images** → Stored in **S3**, delivered with **CloudFront** and responsive image variants.

## 4) Services & Responsibilities

### 4.1 Backend-for-Frontend (BFF)

- Shapes data for the storefront, aggregates calls, enforces caching, rate limits, and A/B flags.
- **Tech:** Node.js (TypeScript), Fastify/Express + OpenAPI.

### 4.2 Catalog Service

- Product, category, attributes, pricing, promotions, inventory summary view.
- **DB:** Aurora Postgres. Write-through to OpenSearch for text/synonym search.

### 4.3 Search Service

- OpenSearch indexes: products, categories, suggestions.
- Features: synonym sets, analyzers, typo tolerance, popularity boosts, personalization hooks.

### 4.4 Cart Service

- Stateless API with carts in Redis (TTL 30 days). Supports anonymous carts (cookie) and merge on login.

## 4.5 Order Service

- Reserves inventory, commits payment intents, creates order records, issues events (OrderCreated/Cancelled/Refunded).
- DB: Aurora (ACID), outbox pattern → Kafka for reliability.

## 4.6 Payments Adapter

- Integrations: Stripe/Adyen/Braintree via tokenization. PCI SAQ-A maintained (no card data on servers).

## 4.7 Inventory Service

- Stock levels, reservations, backorders. Event-driven updates from WMS/ERP if present.

## 4.8 User/Account (UMS)

- Profiles, addresses, preferences, wishlists. PII encrypted at rest. RBAC for admin endpoints.

## 4.9 Recommendations

- Phase 1: Rule-based (best sellers, trending, similar category). Phase 2: ML with user-behavior features.

## 4.10 CMS (Headless)

- Manage home page banners, copy, SEO metadata, feature flags per locale.
- Options: Strapi/Sanity/Contentful.

---

## 5) Data Model (Sketch)

```
-- products
CREATE TABLE product (
    id UUID PRIMARY KEY,
    sku TEXT UNIQUE,
    title TEXT,
    description TEXT,
    price_cents INT,
    currency CHAR(3),
    category_id UUID,
    attributes JSONB,
    images JSONB,
    created_at TIMESTAMPTZ DEFAULT now(),
    updated_at TIMESTAMPTZ DEFAULT now()
);

-- orders
```

```

CREATE TABLE orders (
    id UUID PRIMARY KEY,
    user_id UUID NULL,
    status TEXT CHECK (status IN
    ('pending', 'paid', 'shipped', 'cancelled', 'refunded')),
    total_cents INT,
    currency CHAR(3),
    payment_intent_id TEXT,
    shipping_addr JSONB,
    items JSONB,
    created_at TIMESTAMPTZ DEFAULT now()
);

```

OpenSearch `products` index fields: `title`, `keywords`, `category`, `price`, `popularity`,  
`in_stock`, `attributes.*`.

---

## 6) API Surface (Selected)

- `GET /bff/home` → hero, categories, featured, recommended.
- `GET /search?q=&filters=&sort=&page=`
- `GET /catalog/products/{id}`
- `POST /cart/{cartId}/items` | `DELETE /cart/{cartId}/items/{itemId}`
- `POST /orders` → creates order (idempotency keys)
- `POST /payments/intents` → PSP client secret
- `GET /users/me` | `PUT /users/me`

All APIs documented with **OpenAPI 3.1**, generated SDK for the frontend.

---

## 7) Frontend (Production)

- **Framework:** Next.js 14 (SSR/SSG + ISR) for SEO (maps perfectly to the static HTML you supplied).
  - **Styling:** Tailwind + component library (Radix/HeadlessUI). Your existing CSS can be preserved in global.css.
  - **State:** React Query for data caching; feature flags (Unleash/ConfigCat) for experiments.
  - **i18n & locales:** next-intl; currencies with Dinero.js.
  - **Images:** next/image with CloudFront loader; WebP/AVIF responsive.
  - **Accessibility:** Lighthouse ≥ 95; semantic HTML, focus traps, color contrast.
- 

## 8) Caching Strategy

- **Edge (CloudFront):**
- Cache HTML for SSG/ISR pages (TTL 60–300s) with `stale-while-revalidate`.

- Cache JSON API responses that are read-mostly: catalog/search (TTL 30–120s); bypass for personalized endpoints.
  - Cache product images, CSS, JS with immutable hashes (1 year).
  - **App Cache:** Redis for cart/session and computed fragments (e.g., home page modules).
  - **Client:** HTTP caching + React Query staleTimes.
- 

## 9) Security & Compliance

- **AuthN/Z:** OIDC (Cognito/Keycloak), JWT access tokens, short TTL refresh tokens, PKCE.
  - **Network:** Private subnets for services; Public only for CDN, API Gateway; Service-to-service via mTLS.
  - **AppSec:** WAF managed rules + custom IP reputation; rate limiting at API GW; OWASP ASVS checks in CI.
  - **Data:** PII encryption (KMS), field-level encryption for addresses, secrets in AWS Secrets Manager.
  - **Payments:** PCI SAQ-A (JS tokenization), no PANs processed or stored.
  - **Privacy:** GDPR/CCPA controls (export/delete data APIs), consent banner, cookie partitioning.
  - **Audit:** CloudTrail, immutable logs in S3 with Object Lock (WORM).
- 

## 10) Observability & SRE

- **SLIs/SLOs:**
  - Homepage p95 TTFB < 400ms
  - Search API p95 < 200ms; error rate < 0.2%
  - Checkout availability  $\geq$  99.95%
  - **Tracing:** OpenTelemetry SDK in FE & BE  $\rightarrow$  X-Ray/Jaeger.
  - **Metrics:** RED/USE metrics per service; autoscaling alarms.
  - **Logs:** Structured JSON; correlation IDs propagated.
  - **Synthetic monitoring:** Canary tests (every 1–5 minutes) for search and checkout.
  - **Runbooks:** For common incidents (cache stampede, PSP outage, index lag).
- 

## 11) CI/CD & Environments

- **Branching:** trunk-based + short-lived PRs.
- **Pipelines (GitHub Actions/GitLab CI):**
- Lint/Type-check/Unit tests  $\rightarrow$  SAST/Dependency scan  $\rightarrow$  Build  $\rightarrow$  SBOM  $\rightarrow$  IaC plan  $\rightarrow$  Ephemeral preview  $\rightarrow$  Canary  $\rightarrow$  Prod.
- **Deployments:**
- **Frontend:** Static export to S3 + CloudFront invalidations; SSR via AWS Lambda@Edge or Next on ECS/Fargate.
- **Services:** Containerized (ECS/Fargate or EKS). Blue/green with automatic rollback on error budget burn.
- **Infra as Code:** Terraform + Terragrunt, drift detection.
- **Secrets:** SOPS/Chamber, rotated automatically.

---

## 12) Data & Analytics

- **Event bus:** Kafka/MSK; outbox pattern from services.
  - **Warehouse:** Redshift or Snowflake; ingestion via Kafka Connect/Firehose.
  - **BI:** Looker/QuickSight dashboards: funnel, AOV, LTV, search zero-results, inventory turns.
  - **ML:** Feature store (SageMaker Feature Store) powering recommendations and search boosts.
- 

## 13) Scalability & Resilience

- Multi-AZ for Aurora/OpenSearch; cross-region read replicas for DR.
  - RPO ≤ 5 minutes (binlog/Kafka), RTO ≤ 60 minutes with IaC + backups.
  - Rate limiting and load shedding on BFF; circuit breakers for PSP/Search.
  - Bulkheads: separate autoscaling groups for checkout vs. browse.
- 

## 14) Cost Controls

- Rightsizing + autoscaling; Graviton instances.
  - Tiered storage (S3 IA/Glacier for media versions).
  - OpenSearch hot-warm tiers; index lifecycle policies.
  - FinOps dashboards with cost anomaly detection.
- 

## 15) Migration Path from Current HTML/CSS

1. **Adopt Next.js** and import your HTML/CSS into components/pages; wire header/search/category links to BFF endpoints.
  2. **Introduce BFF** to abstract future microservices; initially proxy to a monolith if needed.
  3. **Stand up Catalog + Search** with minimal schema; ETL seed data.
  4. **Add Cart + Orders + Payments**; connect Stripe test keys; enable SAQ-A.
  5. **Cut over images** to S3 + CloudFront; enable next/image responsive.
  6. **Enable WAF + Auth**; migrate Sign-In/Sign-Up to Cognito hosted UI.
  7. **Observability & SLOs**; add synthetic tests before launch.
  8. **Turn on CI/CD** with preview envs; blue/green deploys.
- 

## 16) Non-Functional Requirements (NFRs)

- **Perf:**
- LCP < 2.5s on 75th percentile 3G/low-end mobile
- CLS < 0.1; TBT < 200ms
- **Security:** No critical/open high CVEs in runtime images; monthly pen-test cadence.
- **Availability:** 99.95% service tier; 99.9% for search.

---

## 17) Example Terraform Modules (outline)

- `modules/network` VPC, subnets, NAT, SGs
  - `modules/ecs_service` service + ALB + autoscaling
  - `modules/rds_aurora` with snapshots/parameter groups
  - `modules/opensearch` domains/ILM
  - `modules/cloudfront_s3` static hosting + OAC + WAF association
  - `modules/msk` Kafka cluster + topics + SCRAM
- 

## 18) Risk Register (initial)

Risk	Impact	Likelihood	Mitigation
OpenSearch cost growth	High	Medium	Hot-warm tiers, shard sizing, ILM
PSP outage	High	Low	Retry w/ backoff, 3DS failover, capture later
Cache stampede	Medium	Medium	Request coalescing, stale-while-revalidate
PII exposure	High	Low	Data minimization, DLP scans, field encryption

---

## 19) Next Steps (Execution-ready Checklist)

- [ ] Create AWS accounts (prod/stage/dev), org SCPs
  - [ ] Stand up VPC + shared services (ECR, MSK, Aurora)
  - [ ] Bootstrap Terraform + CI pipelines
  - [ ] Scaffold Next.js app and BFF with OpenAPI
  - [ ] Implement Catalog + Search MVP
  - [ ] Implement Cart + Orders + Payments (Stripe)
  - [ ] Configure CloudFront/S3 + image pipeline
  - [ ] Enable WAF, Cognito, Secrets Manager
  - [ ] Add telemetry (OTel) + dashboards + SLOs
  - [ ] Run load test, pen test, and chaos drills
- 

**Appendix A:** Sample OpenAPI, IAM policies, and Helm/Compose manifests can be added upon request.